Jinxi(Steven) Yu

Los Angeles, CA | +1 (626) 265 0641 | yujinxi1104@hotmail.com | linkedin.com/in/jinxiy1104

EDUCATION

University of California, Los Angeles (UCLA) — **Los Angeles, CA**
B.S. in ***Computer Science*** *and* ***Applied Mathematics*** — Sep 2024 – Jun 2027
GPA: 3.97/4.00 | **Awards:** Dean's Honors List
Relevant Coursework: Object-Oriented Programming (C++), ML, NLP, DL, Algorithms, Operating System, Data Mining

WORK EXPERIENCE

Yiyue Intelligent Technology Co., Ltd. — **Nanjing, China**
Full-Stack Developer Intern — Jul 2025 – Sep 2025
- Developed key modules (grade & course management, role-based notifications) for a university web using Vue.js and Java/Spring Boot
- Used Spring Data JPA to perform CRUD operations on MySQL without raw SQL, ensuring data integrity across 50+ test cases

SCAI Lab UCLA (AI and Data Mining) — **Los Angeles, CA**
Research Assistant — Feb 2026 – Present
- Designed and implemented an LLM-assisted data mining pipeline to extract 3D molecular structures from ChemRxiv, downloading 47,000+ papers and supplementary files via DOM-based scraping with resumable checkpointing
- Converted unstructured PDFs into structured molecular datasets for downstream ML and quantum chemistry tasks

Sriram Lab UCLA (Computational Systems Biology) — **Los Angeles, CA**
Research Assistant — Oct 2025 – Feb 2026
- Reproduced an end-to-end LLM-based pipeline to predict biological age from structured health data, including preprocessing, prompt-based inference, and downstream statistical analysis
- Validated LLM-predicted bio age against mortality outcomes using Python-based analysis, observing a statistically positive correlation

Huatai Securities Co., Ltd. (Wealth Management Dept.) — **Nanjing, China**
Data Analysis Intern — Jun 2025 – Jul 2025
- Analyzed 2,000+ high-net-worth client records to identify potential investment leads using structured financial metrics
- Cross-referenced foundation and university financial statements to identify 30+ organizations with investable endowment assets

PROJECT HIGHLIGHTS

EgoX - AI-Powered Personal Career Wiki | LangGraph, RAG, Postgres + pgvector, Vercel AI SDK, Drizzle ORM
Sole Developer — Apr 2026 – Present
- Built an AI-native knowledge system with LLM-based structured extraction, vector similarity deduplication, and entity merging with conflict resolution, powering a RAG interface with streaming citations under a $5/month budget
- Designed a hybrid rule + LLM gap-detection Q&A system, using a single structured generation for follow-up questions and feeding responses back into the extraction with scoped updates to prevent hallucination — reducing manual editing to a guided one-click workflow

Mock-Trader | Simulation Platform
Sole Developer — Dec 2025 – Present
- Built a backend-first trading simulation platform with deterministic rule engines and real-time market data, ensuring reproducible evaluation and auditability via append-only trade ledgers
- Implemented a modular Node.js + TypeScript backend with PostgreSQL, including append-only trade ledgers, position reconstruction, and stateless JWT-based authentication

MNIST: HOG-SVM vs ResNet with Interactive Demo | PyTorch
Sole Developer — Feb 2026 – Feb 2026
- Built an end-to-end MNIST pipeline (60k/10k) comparing HOG + LinearSVC and a custom 1-channel ResNet18, with unified preprocessing, evaluation, metrics reporting, and an interactive canvas, allowing users to draw digits and observe real-time model predictions
- Achieved 97.87% test accuracy with HOG (9 orientations, 4x4 cells, 2x2 blocks) and 98.64% with ResNet18 (Adam, batch size 128, 5 epochs, 55k/5k split), a +0.77% accuracy gain

NanoGPT (Transformer Language Model) | PyTorch
Sole Developer — Oct 2025 – Nov 2025
- Reproduced a Transformer-based language model following the NanoGPT architecture, including self-attention and causal masking
- Trained a small-scale autoregressive model to generate Shakespeare-style text, gaining hands-on experience with Transformer

Market Sentiment Predictor | NLP, LLM Embeddings
Team Research Lead — Sep 2025 – Jan 2026
- Built an NLP-based market sentiment classifier using 768-dimensional LLM embeddings on financial news data (50/50 balanced split)
- Evaluated Logistic Regression and a 2-layer MLP, achieving 78.75% accuracy (+3.75% over linear model) and 0.80 F1-score

SKILLS & INTERESTS

Languages: C++, Java, Python, JavaScript, TypeScript, HTML/CSS
Frameworks/Tools: React, Vue, Node.js, PyTorch, Claude Code, Redis, Pandas, Spring Boot, Bash, Git, MySQL, Docker, PostgreSQL